

# An Introduction to GPU and CUDA C/C++

Slides adapted by Dr Sparsh Mittal

Courtesy for slides: NVIDIA, Mutlu/Kirk/Hwu and others

# What is CUDA?

- CUDA Architecture
  - Expose GPU parallelism for general-purpose computing
  - Retain performance
- CUDA C/C++
  - Based on industry-standard C/C++
  - Small set of extensions to enable heterogeneous programming
  - Straightforward APIs to manage devices, memory etc.
- We discuss CUDA C/C++ and GPU architecture (briefly)

# Introduction to CUDA C/C++

- What will you learn in this session?
  - Start from “Hello World!”
  - Write and launch CUDA C/C++ kernels
  - Manage GPU memory
  - Manage communication and synchronization

# Architectural parameters of recent NVIDIA GPUs

RF = register file


	Archite cture	Compute Capability	L1 size (KB)	L2 size (KB)	RF size (KB)	# of SMs	Total RF size (KB)
G80	Tesla	1.0	None	None	32	16	512
GT200	Tesla	1.3	None	None	64	30	1920
GF100	Fermi	2.0	48	768	128	16	2048
GK110	Kepler	3.5	48	1536	256	15	3840
GK210	Kepler	3.7	48	1536	512	15	7680
GM204	Maxwell	5.2	48	2048	256	16	4096
GP100	Pascal	6.0	48	4096	256	56	14336
GV100	Volta	7.0	128	6144	256	80	20480

This PPT applies to devices with capability  $\geq 2.0$

# **HETEROGENEOUS COMPUTING**

# Heterogeneous Computing

- Terminology:
  - *Host* The CPU and its memory (host memory)
  - *Device* The GPU and its memory (device memory)




Host



Device

# GPU vs. CPU



# “The Tradeoff”

Optimizes  
***LATENCY***



**CPU**

Optimizes  
***THROUGHPUT***



**GPU**



# Heterogeneous Computing

```
#include <iostream>
#include <algorithm>

using namespace std;

#define N 1024
#define RADIUS 3
#define BLOCK_SIZE 16

__global__ void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;
    int lindex = threadIdx.x + RADIUS;

    // Read input elements into shared memory
    temp[lindex] = in[gindex];
    if (threadIdx.x < RADIUS) {
        temp[lindex - RADIUS] = in[gindex - RADIUS];
        temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];
    }

    // Synchronize (ensure all the data is available)
    __syncthreads();

    // Apply the stencil
    int result = 0;
    for (int offset = -RADIUS; offset <= RADIUS; offset++)
        result += temp[lindex + offset];

    // Store the result
    out[gindex] = result;
}

void fill_ints(int *x, int n) {
    fill_n(x, n, 1);
}

int main(void) {
    int *in, *out; // host copies of a, b, c
    int *d_in, *d_out; // device copies of a, b, c
    int size = (N + 2 * RADIUS) * sizeof(int);

    // Alloc space for host copies and setup values
    in = (int *)malloc(size); fill_ints(in, N + 2 * RADIUS);
    out = (int *)malloc(size); fill_ints(out, N + 2 * RADIUS);

    // Alloc space for device copies
    cudaMalloc((void **)&d_in, size);
    cudaMalloc((void **)&d_out, size);

    // Copy to device
    cudaMemcpy(d_in, in, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_out, out, size, cudaMemcpyHostToDevice);

    // Launch stencil_1d() kernel on GPU
    stencil_1d<<<N/BLOCK_SIZE, BLOCK_SIZE>>>>(d_in + RADIUS,
    d_out + RADIUS);

    // Copy result back to host
    cudaMemcpy(out, d_out, size, cudaMemcpyDeviceToHost);




    // Cleanup
    free(in); free(out);
    cudaFree(d_in); cudaFree(d_out);
    return 0;
}
```

parallel fn

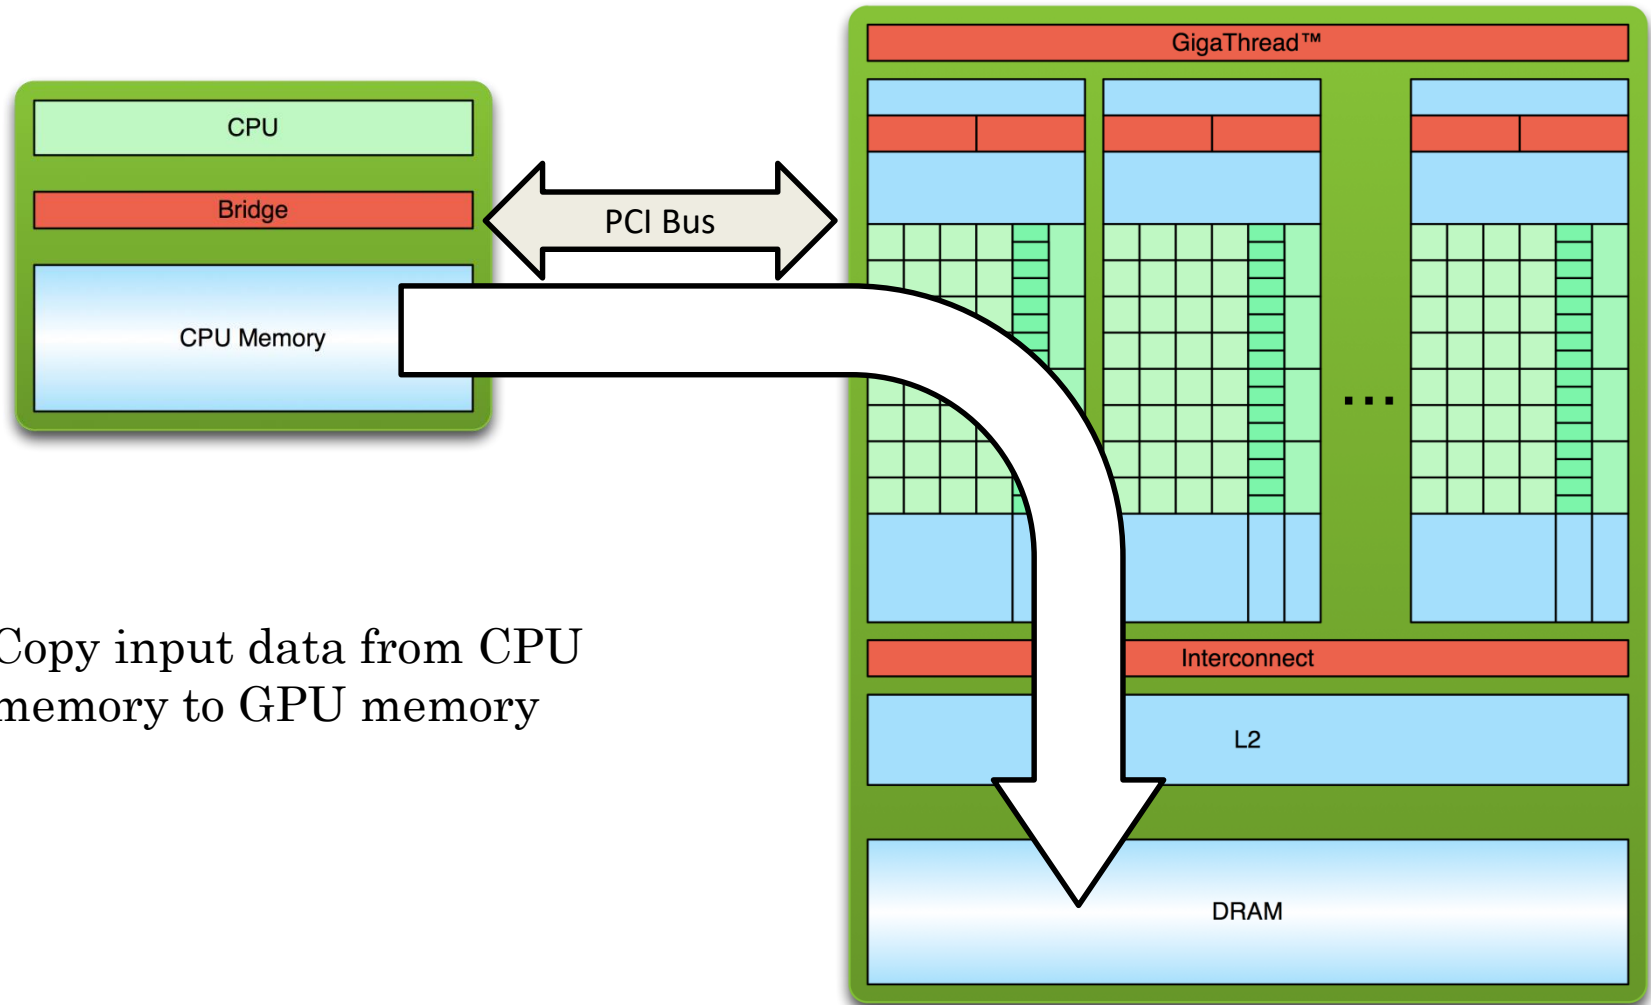
serial code

parallel code

serial code

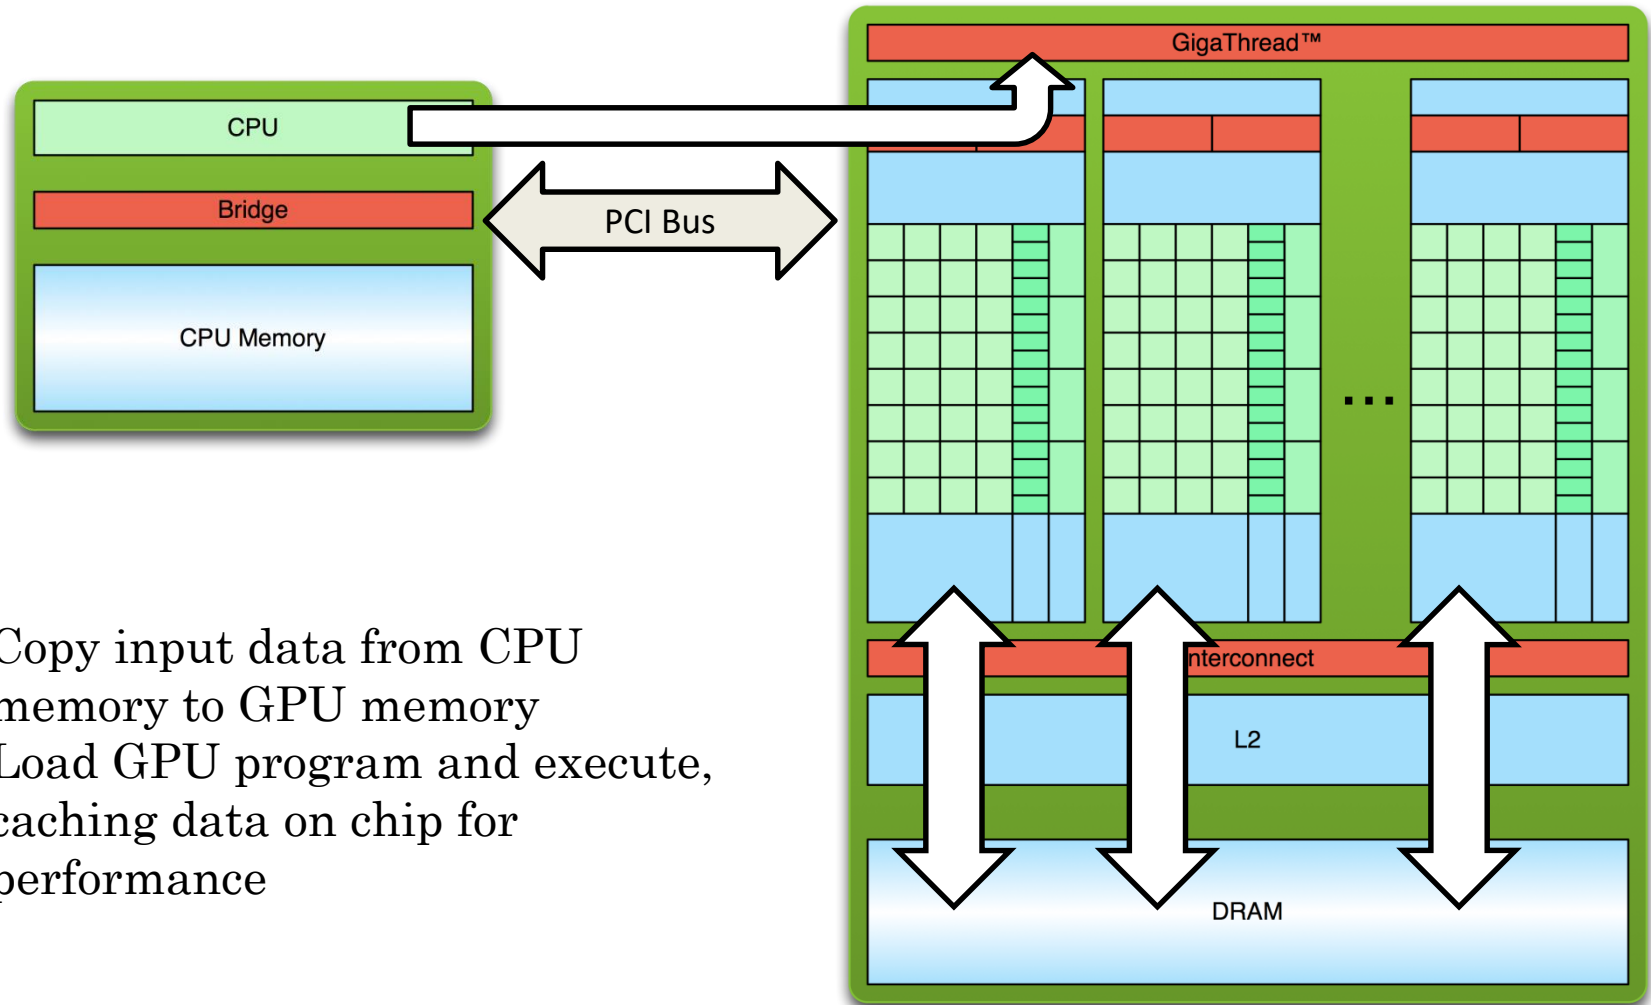


# Simple Processing Flow

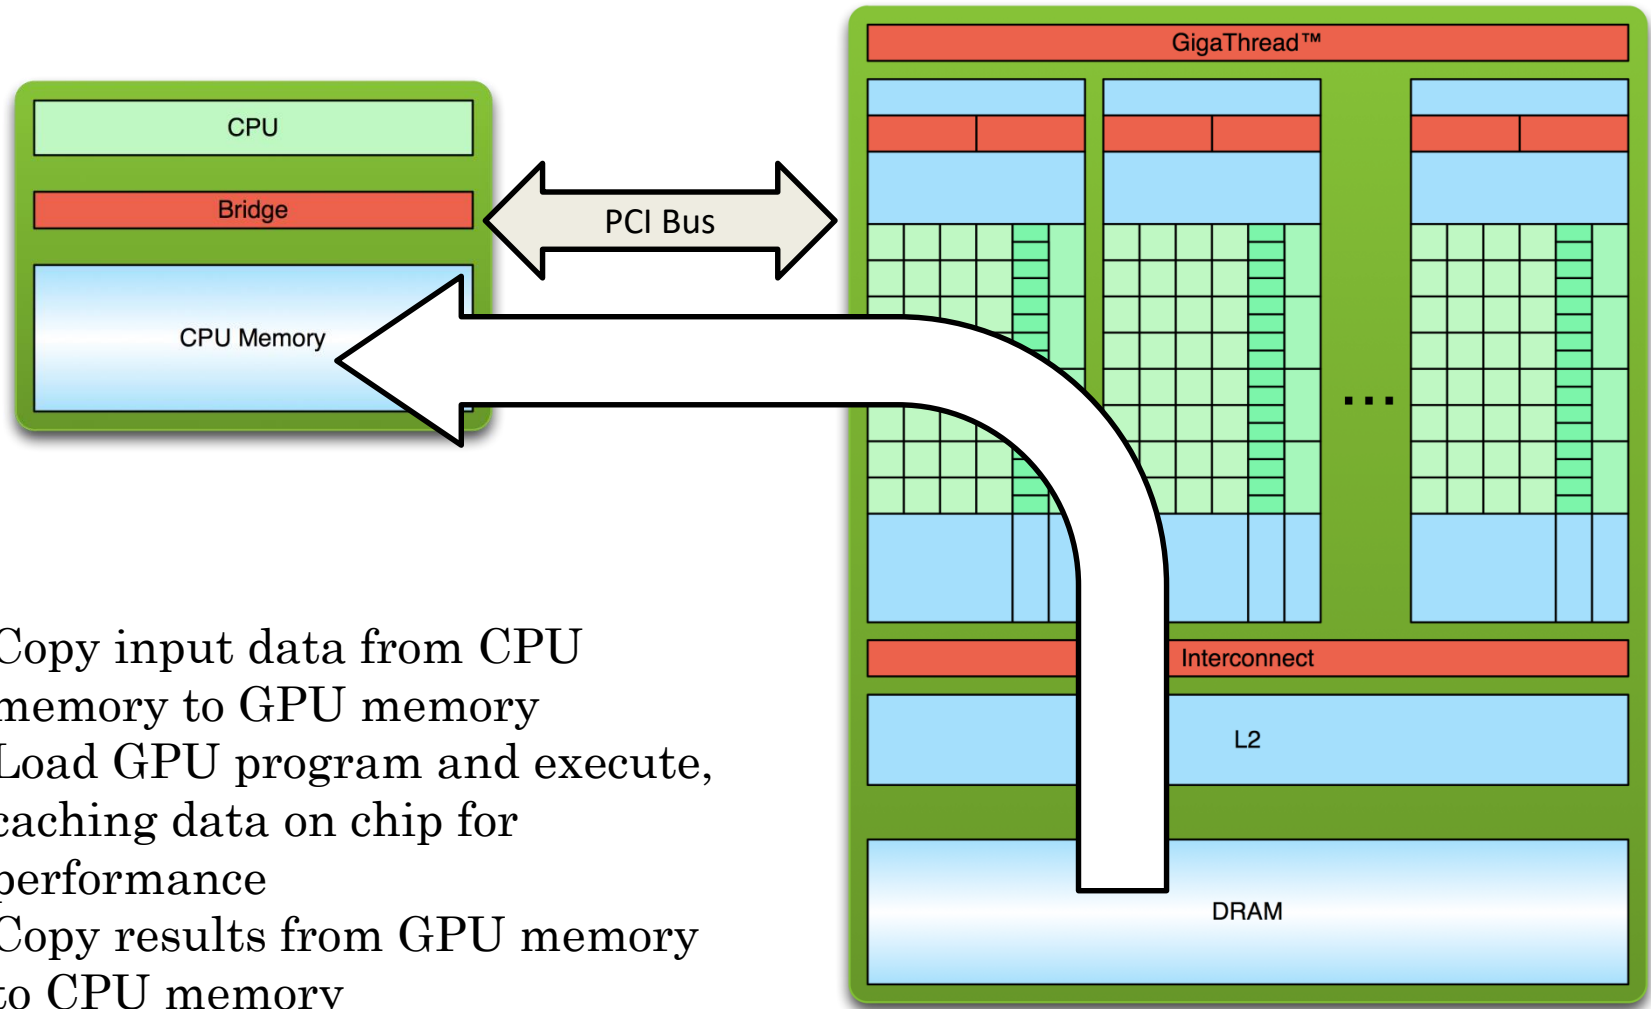


1. Copy input data from CPU memory to GPU memory

# Simple Processing Flow



# Simple Processing Flow



# CUDA extension to declare functions

**\_\_global\_\_** called only from host  
executes only on device

**\_\_device\_\_** called only from device  
executes only on device

**\_\_host\_\_** called only from host  
executes only on host

# Hello World!

```
int main(void) {  
    printf("Hello World!\n");  
    return 0;  
}
```

- Standard C that runs on the host
- NVIDIA compiler (nvcc) can be used to compile programs with no *device* code

Output:

```
$ nvcc  
hello_world.cu  
$ a.out  
Hello World!  
$
```

# Hello World! with Device Code

```
__global__ void mykernel(void) {  
}
```

```
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```

- Two new syntactic elements...

# Hello World! with Device Code

```
__global__ void mykernel(void) {  
}
```

- CUDA keyword `__global__` indicates a function that:
  - Runs on the device
  - Is called from host code
- `nvcc` separates source code into host and device components
  - Device functions (e.g. `mykernel()`) processed by NVIDIA compiler
  - Host functions (e.g. `main()`) processed by standard host compiler
    - `gcc`, `cl.exe`



# Hello World! with Device Code

```
mykernel<<<1,1>>>();
```

- Triple angle brackets mark a call from *host* code to *device* code
  - Also called a “kernel launch”
  - We’ll return to the parameters (1,1) in a moment
- That’s all that is required to execute a function on the GPU!

# Hello World! with Device Code

```
__global__ void mykernel(void) {  
}
```

```
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```

Output:

```
$ nvcc hello.cu  
$ a.out  
Hello World!  
$
```

- In this example, `mykernel()` does nothing

# Printing Hello World from Device

```
//filename helloPrintFromDevice.c
#include <stdio.h>
__device__ const char *STR = "HELLO WORLD!";
const char STR_LENGTH = 12;

__global__ void hello()
{
    printf("%d %c\n", threadIdx.x, STR[threadIdx.x %
STR_LENGTH]);
}
int main(void){
    int num_threads = STR_LENGTH;
    int num_blocks = 1;
    hello<<<num_blocks,num_threads>>>();
    cudaDeviceSynchronize();
    return 0;
}
```

# Output

```
$ nvcc helloPrintFromDevice.cu
```

```
$ ./a.out
```

```
0 H
```

```
1 E
```

```
2 L
```

```
3 L
```

```
4 O
```

```
5
```

```
6 W
```

```
7 O
```

```
8 R
```

```
9 L
```

```
10 D
```

```
11 !
```


```
$
```



Each thread prints one character

# Parallel Programming in CUDA C/C++

- GPU computing is about massive parallelism!
- We will discuss a more interesting example...
- We'll start by adding two integers and build up to vector addition



# Addition on the Device

- A simple kernel to add two integers

```
__global__ void add(int *a, int *b, int *c) {  
    *c = *a + *b;  
}
```

- As before `__global__` is a CUDA C/C++ keyword meaning
  - `add()` will execute on the device
  - `add()` will be called from the host

# Addition on the Device

- Note that we use pointers for the variables

```
__global__ void add(int *a, int *b, int *c) {  
    *c = *a + *b;  
}
```

- `add()` runs on the device, so `a`, `b` and `c` must point to device memory
- We need to allocate memory on the GPU

# Memory Management

- Host and device memory are separate entities
  - *Device* pointers point to GPU memory
    - May be passed to/from host code
    - May *not* be dereferenced in host code
  - *Host* pointers point to CPU memory
    - May be passed to/from device code
    - May *not* be dereferenced in device code
- Simple CUDA API for handling device memory
  - `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`
  - Similar to the C equivalents `malloc()`, `free()`, `memcpy()`





# Addition on the Device: `add()`

- Returning to our `add()` kernel

```
__global__ void add(int *a, int *b, int *c) {  
    *c = *a + *b;  
}
```

- Let's take a look at `main()`...

# Addition on the Device: `main()`

```
int main(void) {  
    int a, b, c; // host copies of a, b, c  
    int *d_a, *d_b, *d_c; // device copies  
    int size = sizeof(int);  
  
// Allocate space for device copies of a, b, c  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
// Setup input values  
    a = 2;  
    b = 7;
```

# Addition on the Device: `main()`

```
// Copy inputs to device
cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);

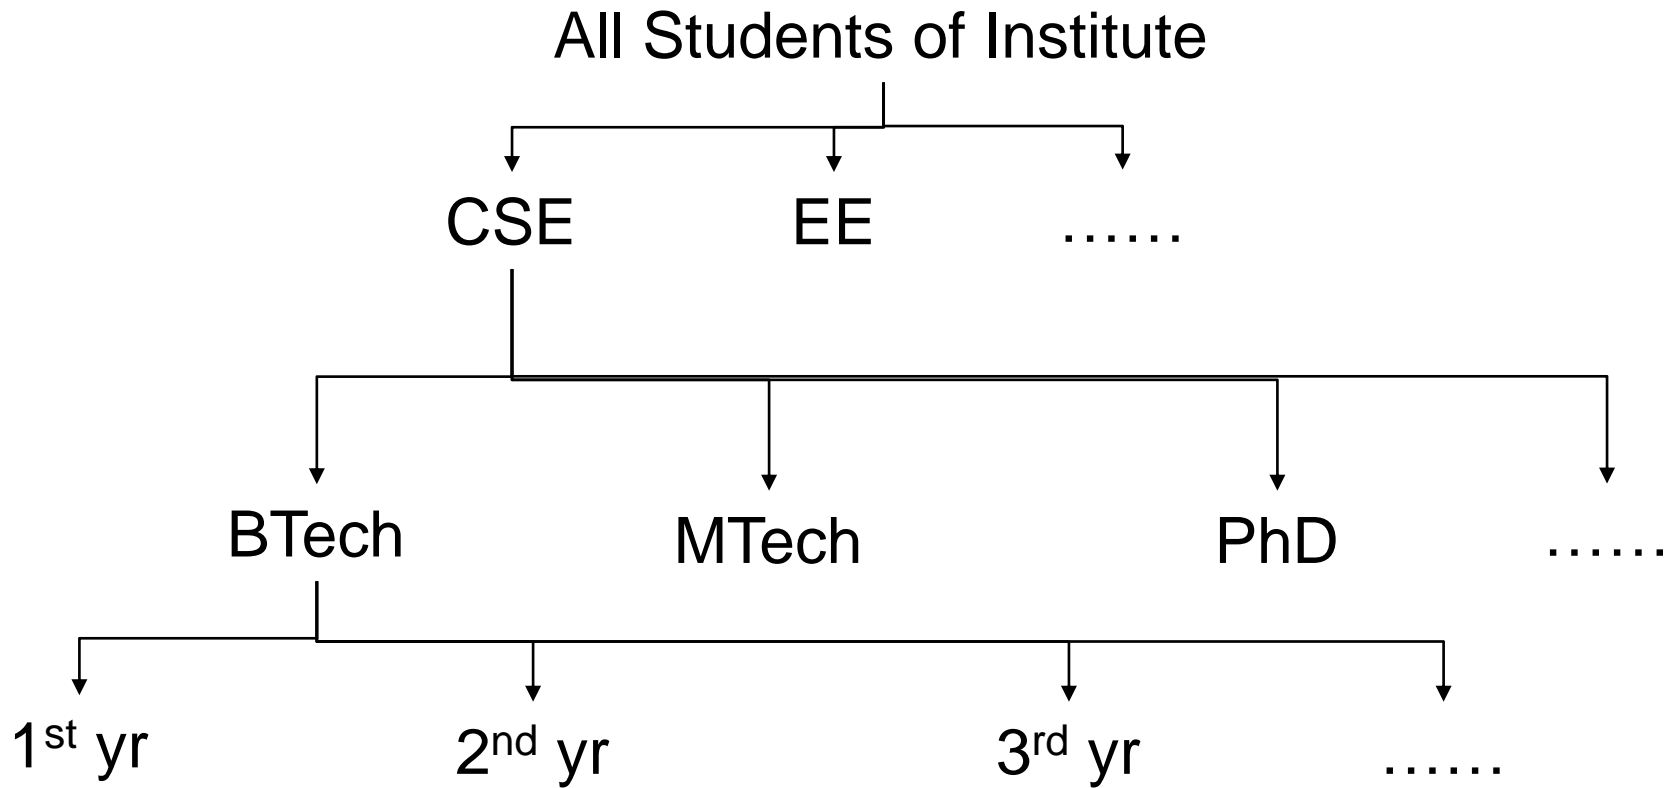
// Launch add() kernel on GPU
    add<<<1,1>>>(d_a, d_b, d_c);

// Copy result back to host
cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);

// Cleanup
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
    return 0;
}
```

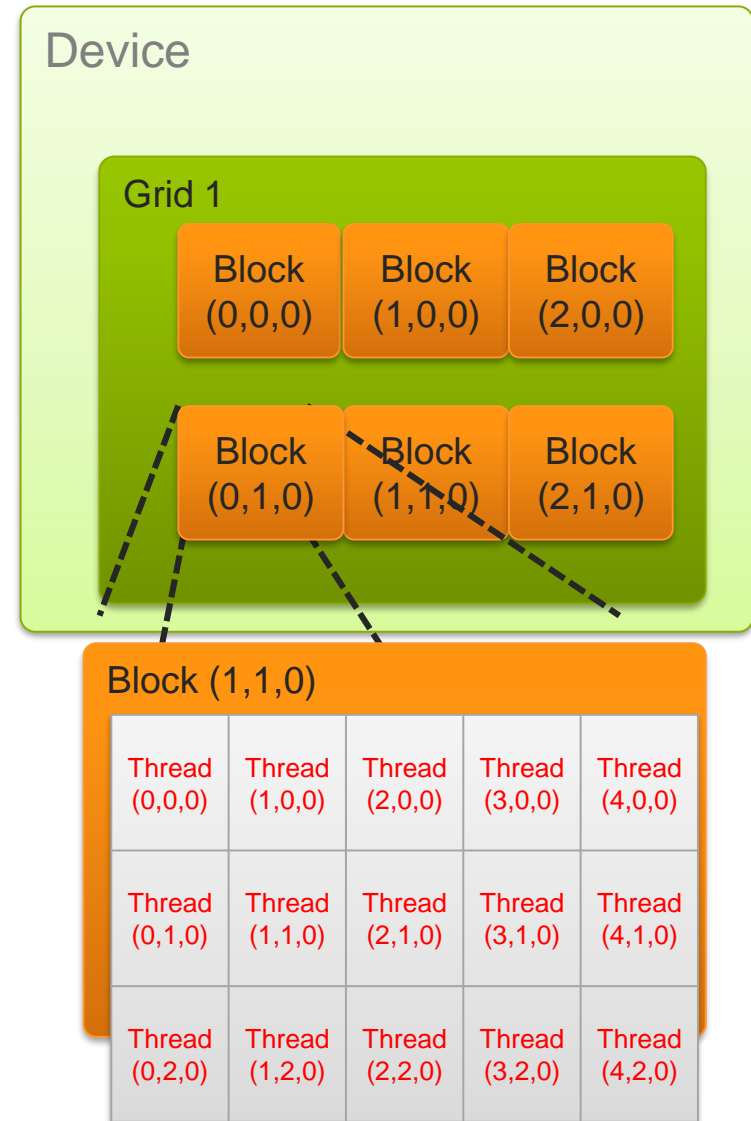
# **UNDERSTANDING THREAD ORGANIZATION**

# Understanding thread organization using example of student groups



# Similarly, threads are organized

- A kernel is launched as a grid of blocks of threads
  - `blockIdx` and `threadIdx` are 3D
  - We showed only one dimension (x)
- Built-in variables:
  - `threadIdx`
  - `blockIdx`
  - `blockDim`
  - `gridDim`




Parallel computing using

**BLOCKS**

# Moving from Scalar to Parallel

- GPU computing is about massive parallelism
  - So how do we run code in parallel on the device?

```
add<<< 1, 1 >>>();
```

```
add<<<  N, 1 >>>();
```

- Instead of executing `add()` once, execute `N` times in parallel



# Vector Addition on the Device

- With **add()** running in parallel we can do vector addition
- Terminology: each parallel invocation of **add()** is referred to as a **block**
  - The set of blocks is referred to as a **grid**
  - Each invocation can refer to its block index using **blockIdx.x**

```
__global__ void add(int *a, int *b, int *c) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x]; }
```

- By using **blockIdx.x** to index into the array, each block handles a different index

# Vector Addition on the Device

```
__global__ void add(int *a, int *b, int *c) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

- On the device, each block can execute in parallel:

Block 0

`c[0]=a[0]+b[0];`

Block 1

`c[1]=a[1]+b[1];`

Block 2

`c[2]=a[2]+b[2];`

Block 3

`c[3]=a[3]+b[3];`

# Vector Addition on the Device:

`add()`

- Returning to our parallelized `add()` kernel

```
__global__ void add(int *a, int *b, int *c)
{
    c[blockIdx.x] = a[blockIdx.x] +
    b[blockIdx.x];
}
```

- Let's take a look at `main()`...

# Vector Addition on the Device: `main()`

```
#define N 512

int main(void) {
    int *a, *b, *c;           // host copies of a, b, c
    int *d_a, *d_b, *d_c;     // device copies
    int size = N * sizeof(int);

    // Alloc space for device copies of a, b, c
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);

    // Alloc space for host copies and initialize
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);
```

# Vector Addition on the Device: `main()`

```
// Copy inputs to device
cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

// Launch add() kernel on GPU with N blocks
    add<<<N,1>>>(d_a, d_b, d_c);

// Copy result back to host
cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);

// Cleanup
    free(a); free(b); free(c);
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
}
```

# Review (1 of 2)

- Difference between *host* and *device*
  - *Host* CPU
  - *Device* GPU
- Using `__global__` to declare a function as device code
  - Executes on the device
  - Called from the host
- Passing parameters from host code to a device function

# Review (2 of 2)

- Basic device memory management
  - `cudaMalloc()`
  - `cudaMemcpy()`
  - `cudaFree()`
- Launching parallel kernels
  - Launch **N** copies of `add()` with `add<<<N,1>>>(...);`
  - Use `blockIdx.x` to access block index

Parallel computing using

**THREADS**



# CUDA Threads

- Terminology: a block can be split into parallel **threads**
- Let's change `add()` to use parallel *threads* instead of parallel *blocks*

```
__global__ void add(int *a, int *b, int *c) {  
    c[threadIdx.x] = a[threadIdx.x] +  
    b[threadIdx.x];  
}
```

- We use **threadIdx.x** instead of **blockIdx.x**
- Need to make one change in **main()**...

# Vector Addition Using Threads: `main()`

```
#define N 512

int main(void) {
    int *a, *b, *c;           // host copies of a, b, c
    int *d_a, *d_b, *d_c;    // device copies of a, b, c
        int size = N * sizeof(int);

// Alloc space for device copies of a, b, c
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);

// Alloc space for host copies and initialize
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);
```

# Vector Addition Using Threads: `main()`

*// Copy inputs to device*

```
cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);  
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
```

*// Launch add() kernel on GPU with N threads*

```
add<<<1,N>>>(d_a, d_b, d_c);
```

*// Copy result back to host*

```
cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);
```

*// Cleanup*

```
free(a); free(b); free(c);
```

```
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
```

```
}
```


# **COMBINING BLOCKS AND THREADS**

# Combining Blocks and Threads

- We've seen parallel vector addition using:
  - Many blocks with one thread each
  - One block with many threads
- Let's adapt vector addition to use both blocks and threads
- Why? We'll come to that...
- First let's discuss data indexing...

# Indexing using Blocks & Threads

- No longer as simple as using `blockIdx.x` and `threadIdx.x`
  - Consider indexing an array with one element per thread (8 threads/block)




- With M threads/block a unique index for each thread is given by:

```
int index = threadIdx.x + blockIdx.x * M;
```

# Indexing Arrays: Example

- Which thread will operate on the red element?



```
int index = threadIdx.x + blockIdx.x * M;  
          =           5   +           2   * 8;  
          = 21;
```

# Vector Addition with Blocks and Threads

- Use the built-in variable `blockDim.x` for threads per block

```
    int index = threadIdx.x + blockIdx.x *  
blockDim.x;
```

- Combined version of `add()` to use parallel threads *and* parallel blocks

```
__global__ void add(int *a, int *b, int *c) {  
    int index = threadIdx.x + blockIdx.x *  
blockDim.x;  
    c[index] = a[index] + b[index]; }
```



# Addition with Blocks and Threads

```
#define N (2048*2048)
#define THREADS_PER_BLOCK 512
int main(void) {
    int *a, *b, *c; // host copies of a, b, c
    int *d_a, *d_b, *d_c; // device copies
    int size = N * sizeof(int);
    // Alloc space for device copies of a, b, c
        cudaMalloc((void **)&d_a, size);
        cudaMalloc((void **)&d_b, size);
        cudaMalloc((void **)&d_c, size);
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);
```

# Addition with Blocks and Threads

*// Copy inputs to device*

```
cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
```

```
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
```

*// Launch add() kernel on GPU*

```
add<<<N/THREADS_PER_BLOCK, THREADS_PER_BLOCK>>>(d_a,  
d_b, d_c);
```

*// Copy result back to host*

```
cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);
```

*// Cleanup*

```
free(a); free(b); free(c);
```

```
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
```

```
}
```

# Handling Arbitrary Vector Sizes

- Typical problems: non-multiples of `blockDim.x`
- Avoid accessing beyond end of the arrays:

```
__global__ void add(int *a, int *b, int *c, int n) {  
    int index = threadIdx.x + blockIdx.x *  
blockDim.x;  
    if (index < n)  
        c[index] = a[index] + b[index];  
}
```

- Update the kernel launch:

```
add<<< (N + M-1) / M, M >>> (d_a, d_b, d_c, N);
```

# Why Bother with Threads?

- Threads seem unnecessary
  - They add a level of complexity
  - What do we gain?
- Unlike parallel blocks, threads have mechanisms to:
  - Communicate
  - Synchronize
- Other constraints: The number of blocks in a single launch is limited (e.g., 65536)
- Number of threads per block is limited

Lets first discuss

# **GPU MEMORY ADDRESS SPACES**

# GPU Memory Address Spaces


1. Local
  2. Shared
  3. Global
- 
- Increasing visibility of data between threads

- In addition there are two more (read-only) address spaces:

1. Constant
2. Texture.

# Local (Private) Address Space


Each thread has its own “local memory”



Note: Location at address 100 for thread 0 is different from location at address 100 for thread 1.

Contains local variables private to a thread.


# Global Address Spaces



- Each thread in the different thread blocks (even from different kernels) can access “global memory”
- **cudaMalloc** allocates global memory
- Threads write their own portion of global memory
- No need for synchronization
- Slow



# Lets take example of Matrix Transpose




# Matrix Transpose

```
__global__ void transpose(float *odata, float* idata, int width, int height){  
    int xIndex = blockIdx.x * TILE_DIM + threadIdx.x;  
    int yIndex = blockIdx.y * TILE_DIM + threadIdx.y;  
  
    int index_in  = xIndex + width * yIndex;  
    int index_out = yIndex + height * xIndex;  
    for (int i=0; i<TILE_DIM; i+=BLOCK_ROWS) {  
        odata[index_out+i] = idata[index_in+i*width];  
    }  
}
```

- “xIndex”, “yIndex”, “index\_in”, “index\_out”, and “i” are in local memory (local variables are register allocated, stack is allocated in local memory)
- “odata” and “idata” are pointers to global memory (both allocated using calls to cudaMalloc -- not shown above)

# Shared Address Space



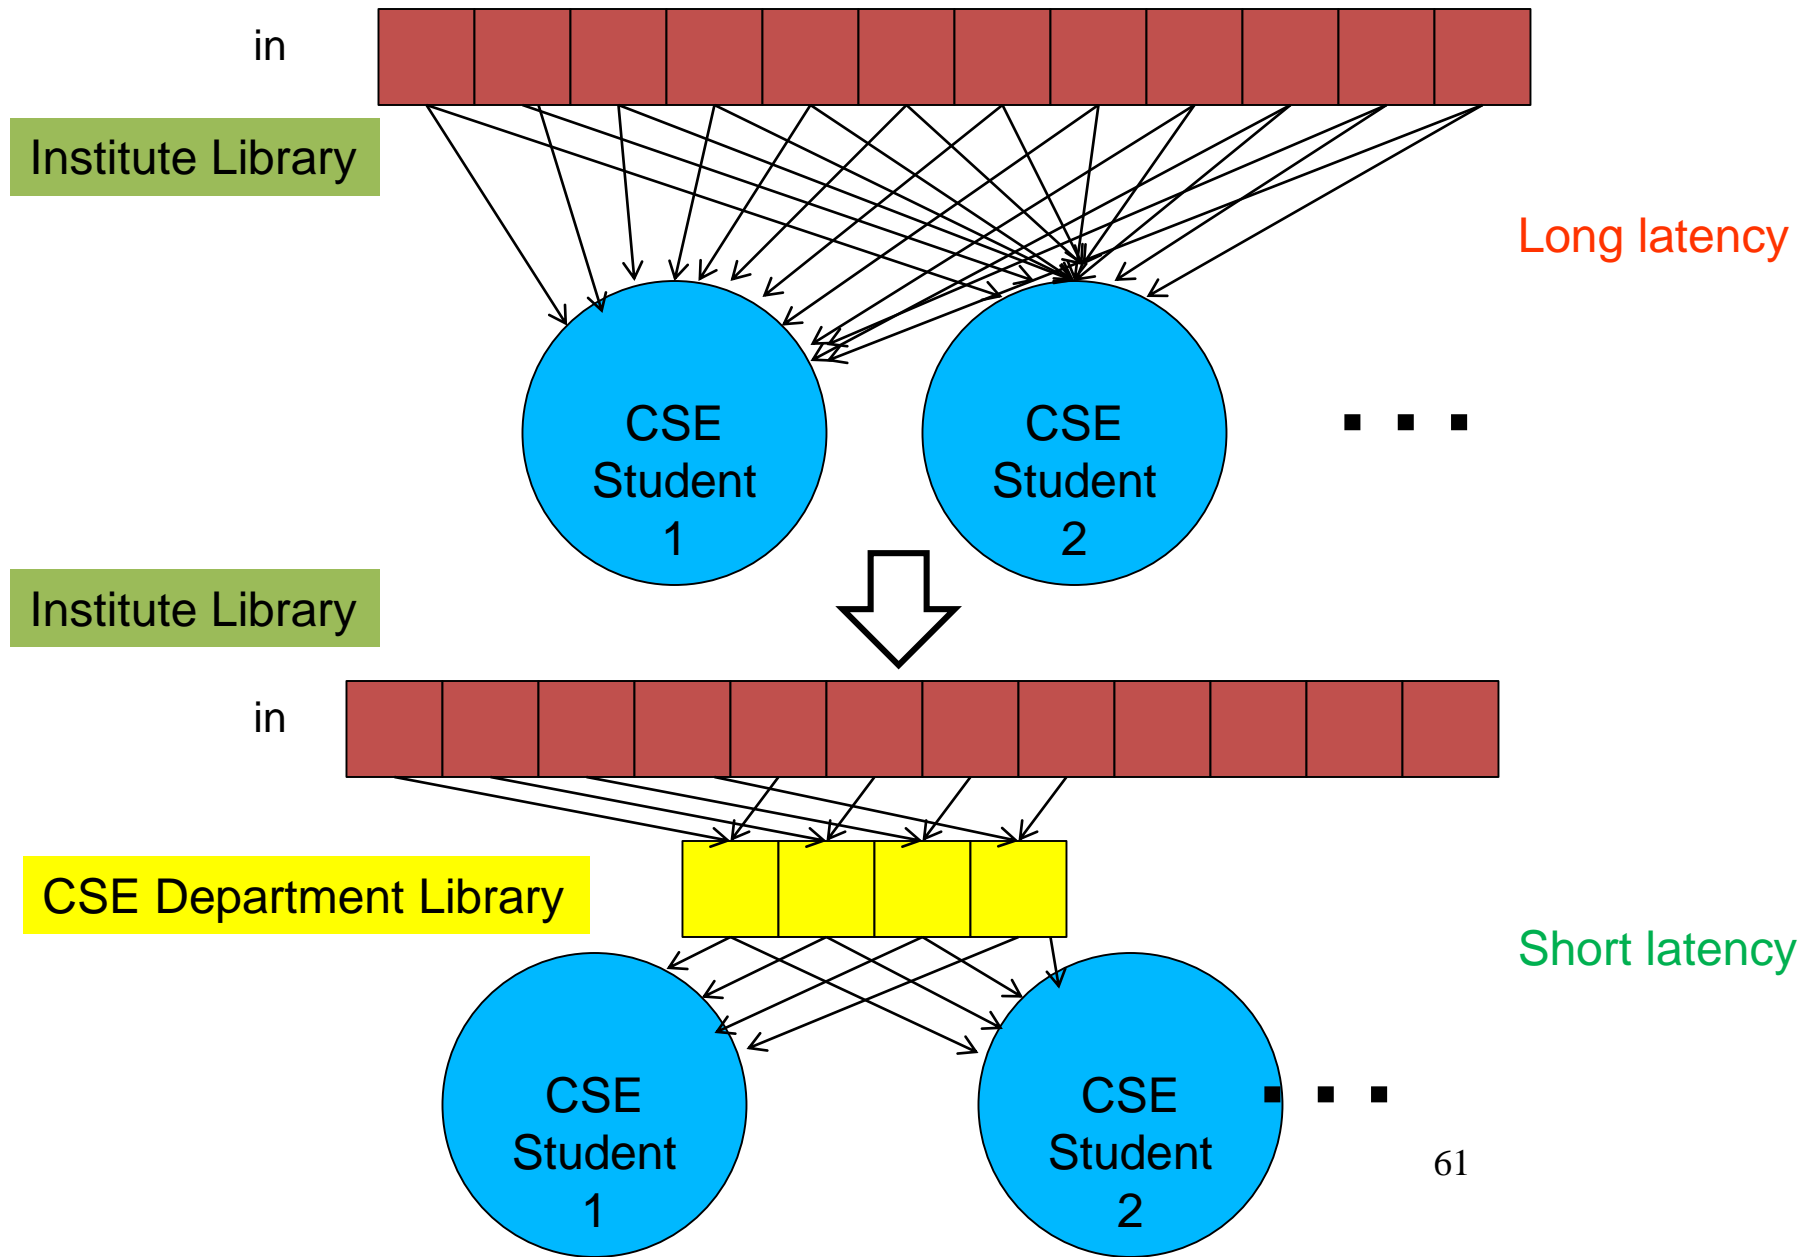
Each thread in the same block can access a memory region called “shared memory”

Limited size (16 to 48 KB).

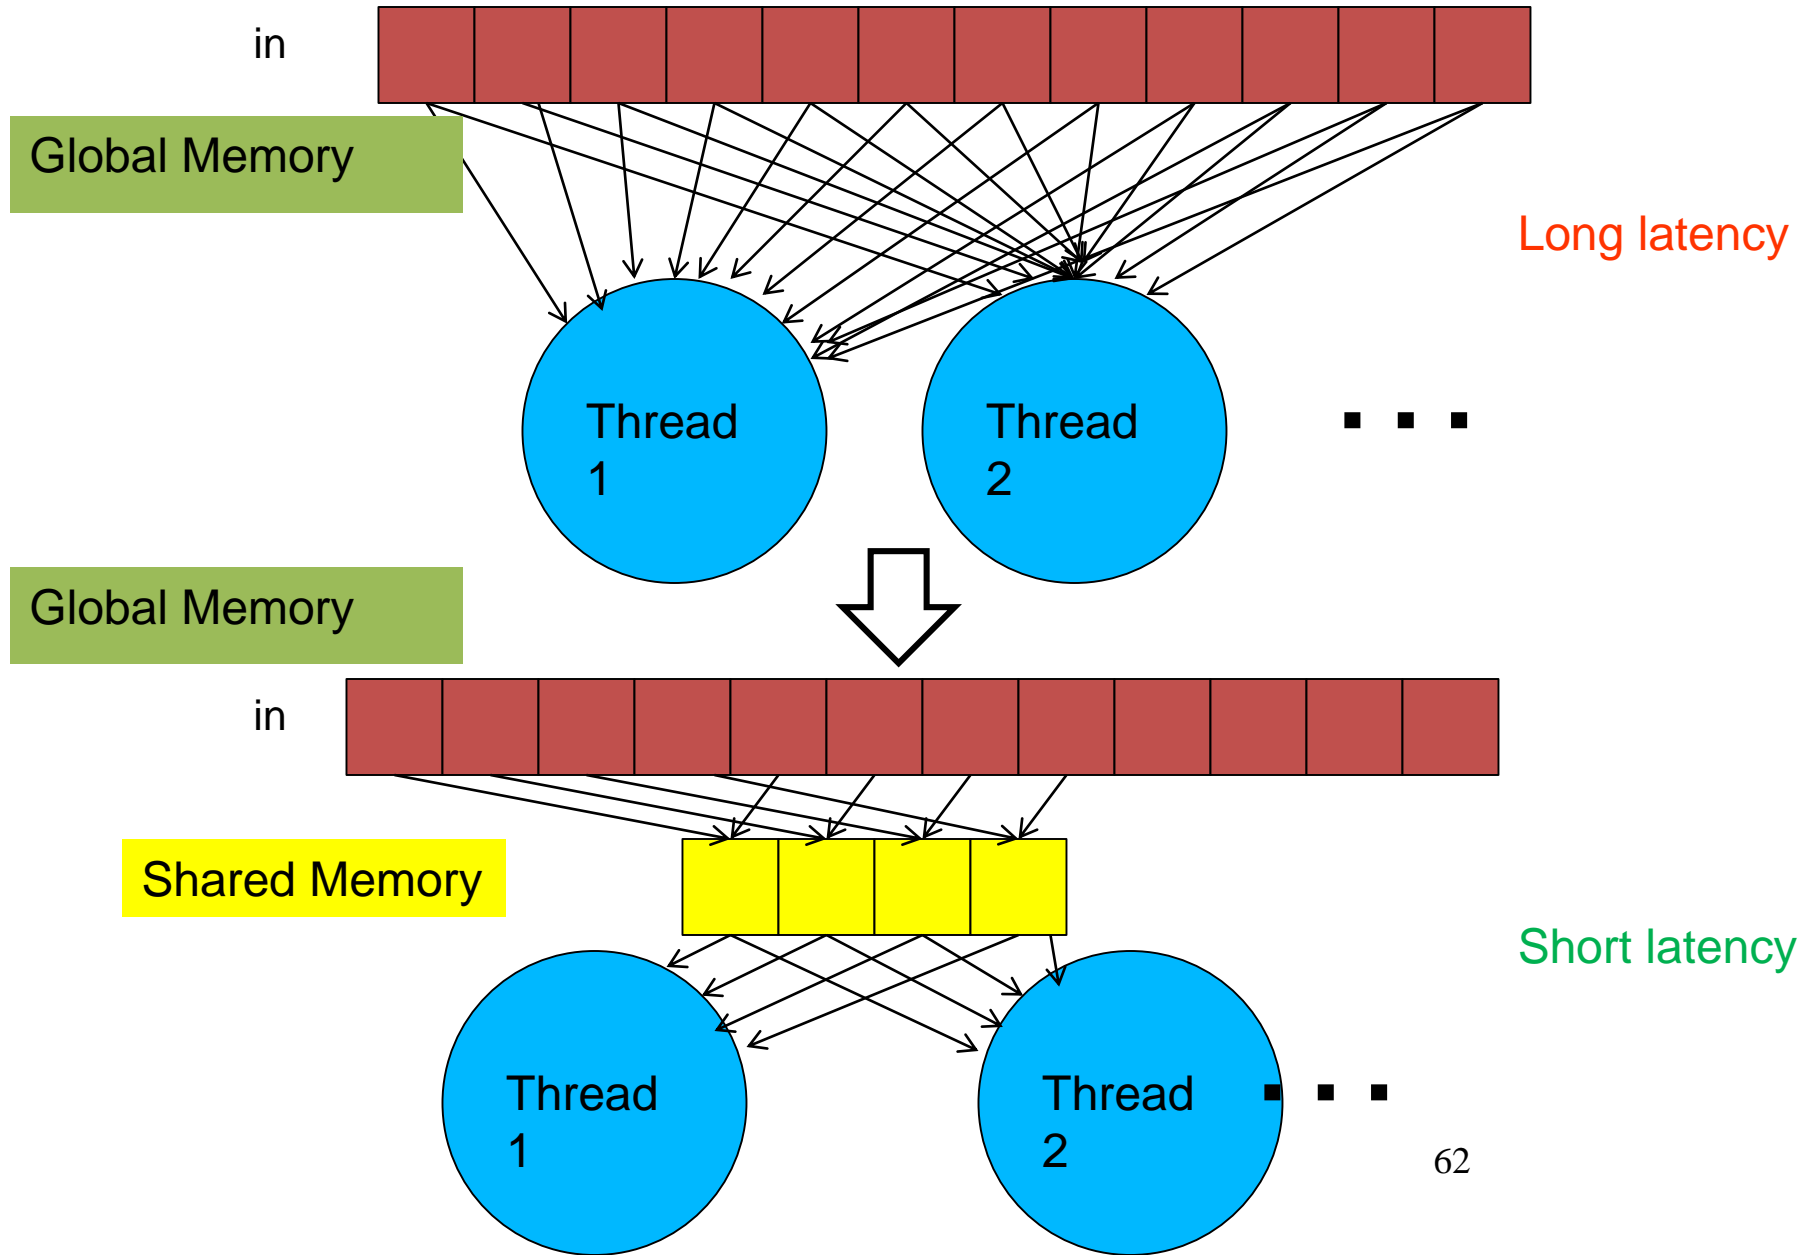
Used as a software managed “cache” to avoid off-chip memory accesses.

Synchronize threads in a thread block using `__syncthreads()`;

# Analogy: Institute and Dept. Library



# Similarly: Global and Shared memory



# CUDA Variable Type Qualifiers

Variable declaration	Memory	Scope	Lifetime	Latency
<code>int LocalVar;</code>	register	thread	thread	1x
<code>int localArray[10];</code>	local	thread	thread	100x
<code>__shared__ int SharedVar;</code>	shared	block	block	1x
<code>__device__ int GlobalVar;</code>	global	grid	application	100x
<code>__constant__ int ConstVar;</code>	constant	grid	application	1x

- **Automatic variables** without any qualifier reside in a **register**
  - **Except per-thread arrays** that reside in local memory
  - Or if there are not enough registers

# Programming scenario 1

## Task:

Load data from global memory

Do **thread-local** computations

Store results to global memory

## Solution:

Load data from global memory

```
float a = d_ptr[blockIdx.x * blockDim.x + threadIdx.x];
```

- Do computation with registers

```
float res = func(a)
```

- Store result

```
d_ptr[blockIdx.x*blockDim.x + threadIdx.x] =res;
```

# Programming scenario 2

**Task:** 1. Load data from global memory 2. Do **block-local** computations 3. Store results to global memory

## **Solution:**

Load data from global memory to shared memory

```
__shared__ float a_sh [ BLOCK_SIZE ];
```

```
int idx = blockIdx .x* blockDim .x + threadIdx .x;
```

```
a_sh [ threadIdx .x] = d_ptr [ idx ];
```

```
__syncthreads ();
```

- Do computation

```
float res = func(a_sh[threadIdx.x])
```

- Store result

```
d_ptr[index] =res;
```




Because it's tricky, lets discuss in more detail:


## **SHARED MEMORY**

# Dept library need synchronization

- Good – when students have similar choices




- Bad – when students have different choices



# Same with Blocking/Tiling

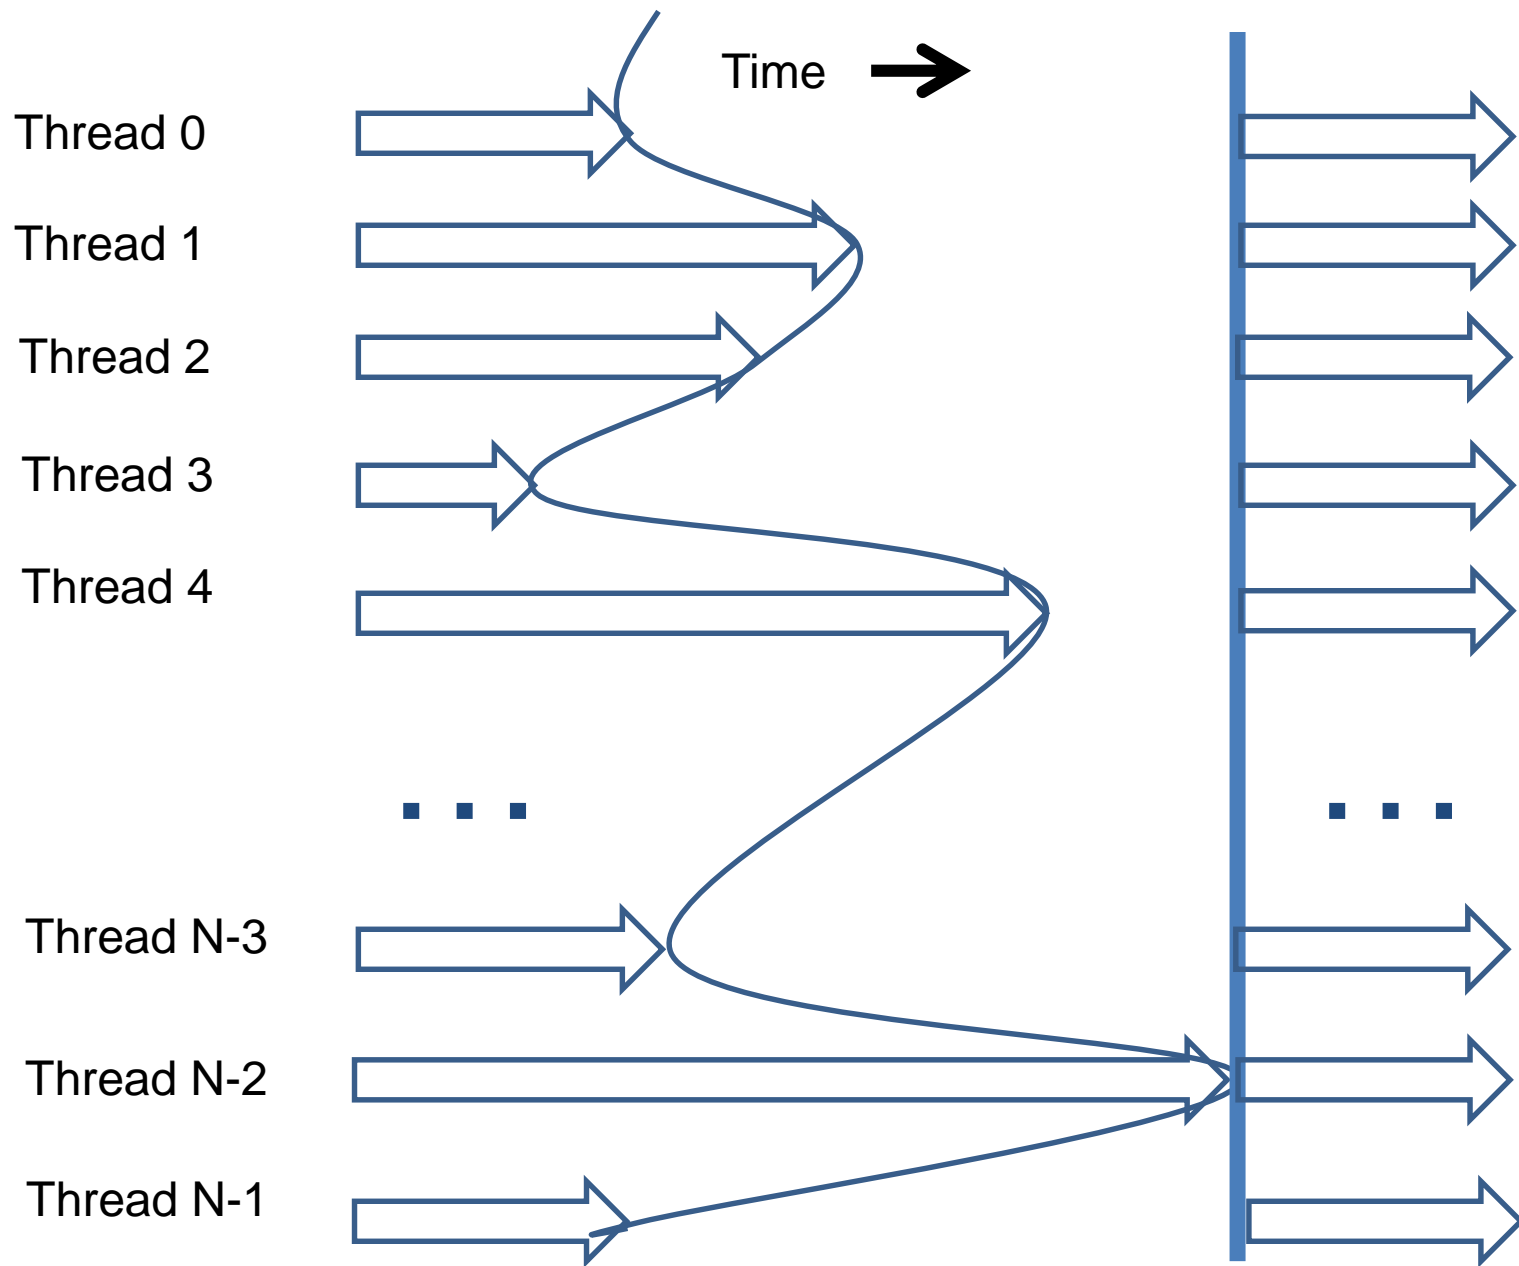
- Good –when threads have similar access timing



- Bad – when threads have very different timing

# Barrier Synchronization

- A function call in CUDA
  - `__syncthreads()`
- All threads in the same block must reach the `__syncthreads()` before any can move on
- Best used to coordinate tiled algorithms
  - To ensure that all elements of a tile are loaded
  - To ensure that all elements of a tile are consumed



An example execution timing of barrier synchronization.

# References

- CUDA language:
  - CUDA by Example, by Jason Sanders and Edward Kandrot, NVIDIA
  - “Programming Massively Parallel Processors: A Hands-on Approach” by David B. Kirk and Wen-mei W. Hwu
- GPU architecture
  - “A Survey of CPU-GPU heterogeneous computing”, S. Mittal et al., CSUR 2015

# Thanks!



Sparsh Mittal [sparsh@iith.ac.in](mailto:sparsh@iith.ac.in)